

Sign-flip vs initial calibration: TTA helps where the model starts poorly calibrated (fill = 95% CI excludes 0)

